

Reinforcement Learning

Advanced Policy Gradient

Александр Костин
telegramm: @Ko3tin
LinkedIn: [kostinalexander](#)

Recap: Policy Gradient

$$\nabla J_{PG}(\theta) = \mathbb{E}_{p_{\theta}(\tau)} \left[\sum_{t=0}^T \nabla \log \pi_{\theta}(A_t | S_t) \sum_{k=t}^T \gamma^{k-t} R_k \right]$$

or

$$\nabla J_{PG}(\theta) = \mathbb{E}_{p_{\theta}(\tau)} \left[\sum_{t=0}^T \nabla \log \pi_{\theta}(A_t | S_t) [Q_{\pi_{\theta}}(S_t, A_t) - b(S_t)] \right]$$

or

$$\nabla J_{AC}(\theta) = \mathbb{E}_{p_{\theta}(\tau)} \left[\sum_{t=0}^T \nabla \log \pi_{\theta}(A_t | S_t) [A_{\pi_{\theta}}(S_t, A_t)] \right]$$

Recap: A2C

- Generate trajectories $\{\tau_i\}$ following $\pi_\theta(a | s)$

- Policy improvement:

Estimate gradient and make gradient ascent step:


$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left[\sum_{t=0}^T \nabla \log \pi_\theta(a_{i,t} | s_{i,t}) A_{\pi_\theta}(s_{i,t}, a_{i,t}) \right]$$

- Policy evaluation:

Estimate gradient and make gradient descent step:

$$\nabla_\phi L(\phi) \approx \frac{1}{N} \sum_{i=1}^N \left[\sum_{t=0}^T \nabla_\phi (r_{i,t} + \gamma \boxed{V_\phi(s_{i,t+1})} - V_\phi(s_{i,t}))^2 \right]$$

Not target network, just frozen parameters



Source

Generalized Advantage Estimation (GAE)

Advantage function estimation:

$$A(s_t, a_t) = Q(s_t, a_t) - V(s_t) = r_t + \gamma V(s_{t+1}) - V(s_t) = -V(s_t) + r_t + \gamma V(s_{t+1})$$

Such advantage estimation might be biased

Also advantage function estimation:

$$A(s_t, a_t) = Q(s_t, a_t) - V(s_t) = -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

Such advantage estimation might have high variance

GAE

$$\hat{A}_t^{(1)} := \delta_t^V = -V(s_t) + r_t + \gamma V(s_{t+1})$$

$$\hat{A}_t^{(2)} := \delta_t^V + \gamma \delta_{t+1}^V = -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2})$$

$$\hat{A}_t^{(3)} := \delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V = -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 V(s_{t+3})$$

$$\hat{A}_t^{(k)} := \sum_{l=0}^{k-1} \gamma^l \delta_{t+l}^V = -V(s_t) + r_t + \gamma r_{t+1} + \cdots + \gamma^{k-1} r_{t+k-1} + \gamma^k V(s_{t+k})$$

$$\hat{A}_t^{(\infty)} = \sum_{l=0}^{\infty} \gamma^l \delta_{t+l}^V = -V(s_t) + \sum_{l=0}^{\infty} \gamma^l r_{t+l}$$

GAE

The generalized advantage estimator $\text{GAE}(\gamma, \lambda)$ is defined as the exponentially-weighted average of these k-step estimators:

$$\begin{aligned}\hat{A}_t^{\text{GAE}(\gamma, \lambda)} &:= (1 - \lambda) \left(\hat{A}_t^{(1)} + \lambda \hat{A}_t^{(2)} + \lambda^2 \hat{A}_t^{(3)} + \dots \right) \\ &= (1 - \lambda) \left(\delta_t^V + \lambda(\delta_t^V + \gamma \delta_{t+1}^V) + \lambda^2(\delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V) + \dots \right) \\ &= (1 - \lambda) \left(\delta_t^V (1 + \lambda + \lambda^2 + \dots) + \gamma \delta_{t+1}^V (\lambda + \lambda^2 + \lambda^3 + \dots) \right. \\ &\quad \left. + \gamma^2 \delta_{t+2}^V (\lambda^2 + \lambda^3 + \lambda^4 + \dots) + \dots \right) \\ &= (1 - \lambda) \left(\delta_t^V \left(\frac{1}{1 - \lambda} \right) + \gamma \delta_{t+1}^V \left(\frac{\lambda}{1 - \lambda} \right) + \gamma^2 \delta_{t+2}^V \left(\frac{\lambda^2}{1 - \lambda} \right) + \dots \right) \\ &= \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}^V\end{aligned}$$

Recap

Policy Gradients and Actor-Critic algorithms are on-policy algorithms so we can not use experience replay. Thus, our sample efficiency is quite low.

Policy Optimisation via Gradient Ascent

Several issues:

- We make gradient step in the space of parameters, get new parameters θ and policy π_θ from θ_{old} and old policy $\pi_{\theta_{old}}$. However, it's difficult to measure the impact of change in parameters on change in policy.
- Apply only first-order optimisation methods
- Low sample efficiency

$$\theta = \theta_{old} + \alpha \nabla J(\theta_{old})$$

Optimisation

$$J(\theta) \approx J(\theta_{old}) + \nabla J(\theta_{old})(\theta - \theta_{old})$$

$$J(\theta) \rightarrow \max_{\theta} \quad \longleftrightarrow \quad \begin{array}{l} \nabla J(\theta_{old})(\theta - \theta_{old}) \rightarrow \max_{\theta} \\ \text{s.t. } (\theta - \theta_{old})^T (\theta - \theta_{old}) \leq \delta \end{array}$$

Let's $d = \theta - \theta_{old}$, then $d^* \propto \nabla J(\theta_{old})$

$$\theta = \theta_{old} + \alpha \nabla J(\theta_{old})$$

Optimisation

$$J(\theta_{old})(\theta - \theta_{old}) \rightarrow \max_{\theta} \text{ s.t.}$$

$$(\theta - \theta_{old})^T K(\theta - \theta_{old}) \leq \delta$$

K is symmetric, positive-definite matrix

Let's $d = \theta - \theta_{old}$, then $d^* \propto K^{-1} \nabla J(\theta_{old})$

$$\theta = \theta_{old} + \alpha K^{-1} \nabla J(\theta_{old})$$

Natural Gradient

$$KL(\pi_{\theta_{old}} || \pi_{\theta}) \approx \frac{1}{2}(\theta - \theta_{old})^T K(\theta_{old})(\theta - \theta_{old}), \text{ where } K(\theta_{old}) = \nabla_{\theta}^2 KL(\pi_{old} || \pi_{\theta})|_{\theta_{old}}$$

$\theta = \theta_{old} + \alpha s$, where $s = K^{-1} \nabla J(\theta_{old})$, α is a step size.

Natural Gradient

$$KL(\pi_{\theta_{old}} || \pi_{\theta}) \approx \frac{1}{2}(\theta - \theta_{old})^T K(\theta_{old})(\theta - \theta_{old}), \text{ where } K(\theta_{old}) = \nabla_{\theta}^2 KL(\pi_{old} || \pi_{\theta})|_{\theta_{old}}$$

$$\theta = \theta_{old} + \alpha s, \text{ where } s = K^{-1} \nabla J(\theta_{old}), \alpha \text{ is a step size.}$$

Choose the largest step:

$$\frac{1}{2}(\theta - \theta_{old})^T K(\theta_{old})(\theta - \theta_{old}) = \delta \iff \alpha = \sqrt{\frac{2\delta}{s^T K s}}$$

$$\theta = \theta_{old} + \alpha K^{-1} \nabla J(\theta_{old})$$

Natural Gradient

$$KL(\pi_{\theta_{old}} || \pi_{\theta}) \approx \frac{1}{2}(\theta - \theta_{old})^T K(\theta_{old})(\theta - \theta_{old}), \text{ where } K(\theta_{old}) = \nabla_{\theta}^2 KL(\pi_{old} || \pi_{\theta})|_{\theta_{old}}$$

$$\theta = \theta_{old} + \alpha s, \text{ where } s = K^{-1} \nabla J(\theta_{old}), \alpha \text{ is a step size.}$$

Choose the largest step:

$$\frac{1}{2}(\theta - \theta_{old})^T K(\theta_{old})(\theta - \theta_{old}) = \delta \iff \alpha = \sqrt{\frac{2\delta}{s^T K s}}$$

$$\theta = \theta_{old} + \alpha K^{-1} \nabla J(\theta_{old})$$

$K \in \mathbb{R}^{|\theta| \times |\theta|}$, K^{-1} computation takes $O(|\theta|^3)$

Conjugate Gradient Method

Paper

K is symmetric, positive-definite matrix


In order to find $K^{-1} \nabla J(\theta_{old})$ we can solve system $Ks = \nabla J(\theta)$ iteratively.

Conjugate Gradient Method

Paper

K is symmetric, positive-definite matrix

In order to find $K^{-1} \nabla J(\theta_{old})$ we can solve system $Ks = \nabla J(\theta)$ iteratively.



Source

Optimisation in Policy Space

Lemma:

$$J(\pi) = J(\pi_{old}) + \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^T \gamma^t A_{\pi_{old}}(S_t, A_t) \right], \text{ where } A_{\pi_{old}}(S_t, A_t) = Q_{\pi_{old}}(S_t, A_t) - V_{\pi_{old}}(S_t)$$

Let's rewrite it as a sum over states instead of timesteps:

$$J(\pi) = J(\pi_{old}) + \sum_s \rho_\pi(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a), \text{ where } \rho_\pi(s) = \mathbb{P}(s_0 = s) + \gamma \mathbb{P}(s_1 = s) + \dots$$

Optimisation in Policy Space

Lemma:

$$J(\pi) = J(\pi_{old}) + \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^T \gamma^t A_{\pi_{old}}(S_t, A_t) \right], \text{ where } A_{\pi_{old}}(S_t, A_t) = Q_{\pi_{old}}(S_t, A_t) - V_{\pi_{old}}(S_t)$$

Let's rewrite it as a sum over states instead of timesteps:

$$J(\pi) = J(\pi_{old}) + \sum_s \rho_\pi(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a), \text{ where } \rho_\pi(s) = \mathbb{P}(s_0 = s) + \gamma \mathbb{P}(s_1 = s) + \dots$$

If this term is nonnegative then the policy improvement is guaranteed

Optimisation in Policy Space

Lemma:

$$J(\pi) = J(\pi_{old}) + \mathbb{E}_{\tau \sim \pi} \left[\sum_{t=0}^T \gamma^t A_{\pi_{old}}(S_t, A_t) \right], \text{ where } A_{\pi_{old}}(S_t, A_t) = Q_{\pi_{old}}(S_t, A_t) - V_{\pi_{old}}(S_t)$$

Let's rewrite in sum over states instead of timesteps:

$$J(\pi) = J(\pi_{old}) + \sum_s \rho_\pi(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a), \text{ where } \rho_\pi(s) = \mathbb{P}(s_0 = s) + \gamma \mathbb{P}(s_1 = s) + \dots$$

If this term is nonnegative then the policy improvement is guaranteed

Since we don't know π this expression is intractable...

Optimisation in Policy Space

$$J(\pi) = J(\pi_{old}) + \sum_s \rho_{\pi}(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a)$$

$$J(\pi) \approx J(\pi_{old}) + \sum_s \rho_{\pi_{old}}(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a) = L_{\pi_{old}}(\pi)$$

Optimisation in Policy Space

$$J(\pi) = J(\pi_{old}) + \sum_s \rho_{\pi}(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a)$$

$$J(\pi) \approx J(\pi_{old}) + \sum_s \rho_{\pi_{old}}(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a) = L_{\pi_{old}}(\pi)$$

If π_{θ} is quite close to $\pi_{\theta_{old}}$ ($\mathbb{E}_{s \sim \rho_{old}}[KL(\pi_{\theta_{old}} || \pi_{\theta})] \leq \delta$), then

$$L_{\pi_{\theta_{old}}}(\pi_{\theta_{old}}) = J(\pi_{\theta_{old}})$$


$$\nabla_{\theta} L_{\pi_{\theta_{old}}}(\pi_{\theta})|_{\theta_{old}} = \nabla_{\theta} J(\pi_{\theta})|_{\theta_{old}}$$

Optimisation in Policy Space

$$J(\pi) = J(\pi_{old}) + \sum_s \rho_{\pi}(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a)$$

$$\begin{aligned} J(\pi) &\approx J(\pi_{old}) + \sum_s \rho_{\pi_{old}}(s) \sum_a \pi(a | s) A_{\pi_{old}}(s, a) = \\ &= J(\pi_{old}) + \sum_s \rho_{\pi_{old}}(s) \sum_a \pi_{old}(a | s) \frac{\pi(a | s)}{\pi_{old}(a | s)} A_{\pi_{old}}(s, a) \\ &= J(\pi_{old}) + \mathbb{E}_{\rho_{old}} \left[\frac{\pi(a | s)}{\pi_{old}(a | s)} A_{\pi_{old}}(s, a) \right] \end{aligned}$$

Visualisation



We want to compute loss function here!

[Source](#)

Trust Region Policy Optimisation (TRPO)

[Original paper](#)

We have to solve the following optimisation problem to generate a policy update:

$$\begin{aligned} \max_{\theta} \quad & \mathbb{E}_{s \sim \rho_{old}, a \sim \pi_{\theta_{old}}} \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_{old}}(a | s)} A_{\pi_{\theta_{old}}}(s, a) \right] \\ \text{s.t.} \quad & \mathbb{E}_{s \sim \rho_{old}} [KL(\pi_{\theta_{old}} || \pi_{\theta})] \leq \delta \end{aligned}$$

The authors change the advantage function by the Q -function.

TRPO Algorithm

Repeat until convergence:

1. Collect transitions following current policy $\pi_{\theta_{old}}$

2. Compute $g = \nabla_{\theta} \frac{1}{N} \sum_{i=1}^N \frac{\pi_{\theta}(a_i | s_i)}{\pi_{\theta_{old}}(a_i | s_i)} Q_{\pi_{\theta_{old}}}(s_i, a_i)$

3. Compute $K = \nabla_{\theta}^2 \frac{1}{N} \sum_{i=1}^N KL(\pi_{\theta_{old}}(. | s_i) || \pi_{\theta}(. | s_i))$

4. Find optimal direction via Conjugate Gradients Method (find $s = K^{-1}g$)

5. Do linear search in optimal direction checking the KL constraint and

objective value for each new parameter: $\theta_j = \theta_{old} + \alpha_j \sqrt{\frac{2\delta}{g^T s}} s$

TRPO

- + Extremely stable
- + Prominent results
- Computational expensive
- Require cheap sampling
- Difficult to implement

Conditional vs Unconditional Problem

TRPO problem

$$\begin{aligned} \max_{\theta} \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_{old}}(a | s)} \hat{A}_t \right] \\ \text{s.t. } \hat{\mathbb{E}}_t [KL(\pi_{\theta_{old}} || \pi_{\theta})] \leq \delta \end{aligned}$$

Equivalent problem

$$\max_{\theta} \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_{old}}(a | s)} \hat{A}_t - \beta KL(\pi_{\theta_{old}} || \pi_{\theta}) \right]$$

\hat{A} is an estimator of the advantage function at timestep t . Here, the expectation $\hat{\mathbb{E}}_t$ indicates the empirical average over a finite batch of samples, in an algorithm that alternates between sampling and optimization.

PPO Objective


$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{old}}(a_t | s_t)}, \text{ so } r_t(\theta_{old}) = 1$$

$$L^{CPI}(\theta) = \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a | s)}{\pi_{\theta_{old}}(a | s)} \hat{A}_t \right] = \hat{\mathbb{E}}_t [r_t(\theta) \hat{A}_t]$$


$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min (r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$

PPO Objective

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min (r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$



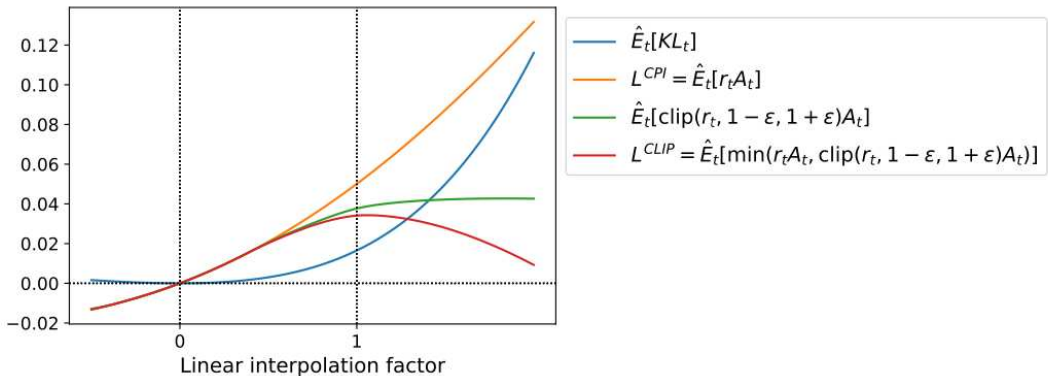
[Source](#)



$p_t(\theta) > 0$	A_t	Return Value of \min	Objective is Clipped	Sign of Objective	Gradient
$p_t(\theta) \in [1 - \epsilon, 1 + \epsilon]$	+	$p_t(\theta) A_t$	no	+	✓
$p_t(\theta) \in [1 - \epsilon, 1 + \epsilon]$	-	$p_t(\theta) A_t$	no	-	✓
$p_t(\theta) < 1 - \epsilon$	+	$p_t(\theta) A_t$	no	+	✓
$p_t(\theta) < 1 - \epsilon$	-	$(1 - \epsilon) A_t$	yes	-	0
$p_t(\theta) > 1 + \epsilon$	+	$(1 + \epsilon) A_t$	yes	+	0
$p_t(\theta) > 1 + \epsilon$	-	$p_t(\theta) A_t$	no	-	✓

[Source](#)

Surrogate Objectives



Source

Optimisation

Like in A2C, we obtain the following object, which is approximately maximised each iteration:

$$L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 H_{\pi_\theta}(s_t),$$

where c_1, c_2 are coefficients, H denotes an entropy bonus,

$L_t^{VF}(\theta)$ is a squared-error loss $\hat{\mathbb{E}}_t[V_\theta(s_t) - V_t^{target}]$

TRPO vs PPO


- Works for smaller models
- + Second-order optimisation

- + Works for big models
- First-order optimisation

Continuous action spaces

- We can learn critic easily
- The problem is finding

$$a_{opt} = \arg \max_a Q(s, a)$$



DDPG

Idea: learn a separate network to find a_{opt}

- Train critic $Q(s,a)$

$$\underset{\theta}{\operatorname{argmin}} \left(Q(s_t, a_t) - [r + \gamma \cdot V(s_{t+1})] \right)^2$$

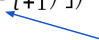
- Train actor $a_{\text{opt}}(s) = \mu(s)$

$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$

DDPG

Idea: learn a separate network to find a_{opt}

- Train critic $Q(s,a)$

$$\underset{\theta}{\operatorname{argmin}} \left(Q(s_t, a_t) - [r + \gamma \cdot V(s_{t+1})] \right)^2$$


- Train actor $a_{\text{opt}}(s) = \mu(s)$

How to get $V(s')$?

$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$

DDPG

Idea: learn a separate network to find a_{opt}


- Train critic $Q(s,a)$

$$\underset{\theta}{\operatorname{argmin}} \left(Q(s_t, a_t) - [r + \gamma \cdot Q(s_{t+1}, \mu_{\theta}(s_{t+1}))] \right)^2$$

- Train actor $a_{\text{opt}}(s) = \mu(s)$

$$\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$$

DDPG



DDPG

Gradient approximation: $\nabla_{\theta} J = \frac{\partial Q^{\theta}(s, a)}{\partial a} \frac{\partial \mu(s|\theta)}{\partial \theta}$

